

# Exam Solutions

## 1. Interface Stability

This was discussed in reading section(s) Interface Stability

This was discussed in lecture section(s) 2C,10F

- a. consequences of an incompatible ABI change

Application programs purchased/obtained by the customer might stop working after the customer upgraded to the new OS version.

- b. how those problems would be responded to

The Independent Software Vendor would have to discover the new interfaces, modify the program to work with the new interfaces, rebuild it, and get the new version out to affected customers. They might also have to distribute different versions of their software to run on different versions of the OS.

- c. how clear interface specifications, distinct from implementation, might help

If the interface specifications were well abstracted from the current implementation this reduces the likelihood that future implementation changes would necessitate incompatible API/ABI changes.

Additionally having a clear written interface specification would raise the visibility of the interface, perhaps making it more obvious to the OS supplier that they were making a change to a committed interface, and that in doing so they were likely to break existing third party applications. If the interfaces were not clearly documented, people would be less likely to consider changes to be important.

## 2. Resource Convoys

This was discussed in reading section(s) ~A7.3

This was discussed in lecture section(s) 7K

- a. A resource convoy is a persistent queue of processes waiting to get access to a popular resource, which eliminates parallelism, increases delays and reduces system throughput.
- b. The key to convoy formation is that processes are no longer able to immediately allocate the required resource, but are always forced to block (until the resource is

freed by the current owner and other processes in line run). Once this happens, the mean service time can easily exceed the mean inter-request time, and the line becomes permanent. It may be precipitated by a process becoming blocked or preempted while holding the resource.

c. Techniques for reducing contention include:

- eliminate mutual exclusion by making the resource truly sharable.
- reduce mutual exclusion by implementing read/write locks.
- reduce contention by breaking up the one resource into a number of sub-resources.
- reduce likelihood of conflict by shortening the protected critical section, or using it less often.
- reduce the likelihood of preemption by moving potentially blocking operations out of the critical section,

### 3. Causes of blocking/preemption

This was discussed in reading section(s) AD4.4,A7.5-7

This was discussed in lecture section(s) 3F,4C

- a. A running process might be preempted if its time slice ends, if its priority drops, or if a higher priority process becomes runnable.
- b. A running process might become blocked if it requests a resource that is not immediately available, or I/O operation. Also, it is (in some sense) blocked when it is swapped out ... since it cannot run until it is swapped back in.

### 4. Evaluating Mutual Exclusion

This was discussed in reading section(s) AD28

This was discussed in lecture section(s) 7D,E

- a. The text identified the key criteria as successful *mutual exclusion*, *fairness* (vs starvation) and *performance* (single processor, multi-processor). I added to this *progress*, not blocking for an available resource and likelihood of avoiding convoys and deadlocks.
- b. Spin locks work (modulo interrupt), and are prone to starvation. They are likely to be quite wasteful if there is contention, but can be very efficient for uncontended use. But they score well on the progress criterion.
- c. Interrupt disables are not usable from user mode and are ineffective against multi-processor parallelism. They are relatively expensive operations, but

relatively fair.

- d. Mutexes guarantee mutual exclusion. Mutexes work to ensure mutual exclusion. They are (with queuing) relatively fair, but there is a race condition where a new locker can get the mutex before the awakened guy at the front of the queue can do so. But this satisfies the progress criterion. The system call, as well as blocking and dispatching are all relatively expensive operations, but blocking is usually much more efficient than spinning.

## 5. DLLs vs Shared Libraries

This was discussed in reading section(s) Linking & Libs

This was discussed in lecture section(s) 3Y

- a. The major capabilities that come with DLLs are
- the ability to open and load (at run-time) modules that did not exist at link time,
  - deferring loading until the modules are actually called,
  - the ability to perform per-module initialization and shut-down
  - the ability to resolve references from the loaded module back into the main program.
- b. Examples of the exploitation of each capability are
- explicit selection and loading is exploited by browser plug-ins which can be obtained long after the browser
  - deferred binding can significantly improve performance (by reducing work at initial program load time) and make it possible for a program to get the benefits of modules that become available after the program starts. This can have a significant performance impact if many plug-ins might be used, but actual use is few and seldom.
  - per module initialization could be used to allocate and initialize private data, register instances, and other complex startup (or shut-down). Device drivers, for instance, require both.
  - the ability to make calls back into the containing program is important if it provides rich services for the plug-in. Here, again, device drivers (which make heavy use of DKI services) are a very good example.
- c. The big extra mechanism that DLLs require is a run-time loader. Why? Because they have to be loaded at run time! They also require a linkage editor that is capable of generating Procedure Linkage Table entries ... but this is a much simpler thing.

## 6. Messages vs shm IPC

This was discussed in reading section(s) mmap(2),send(2),recv(2)

This was discussed in lecture section(s) 7A

- a. The primary advantage of shared memory over message IPC is performance.
- b. Shared memory IPC allows large amounts of data can be transferred, at memory speed, with ordinary user-mode instructions, without the need to make expensive calls to operating system.
- c. The biggest advantage of messages is that they can easily be sent to processes on other machines, whereas shared memory can only be used between processes on a single machine (it can be turned into messages, but doing so sacrifices its performance advantages). This gives us much greater flexibility in how we structure our applications and systems.

Messages sent through the operating system can have authenticated sender identity, and the OS can ensure the integrity and privacy of the message contents. This is because the messages are buffered in, and delivered by the OS ... which does not happen with shared memory.

Also options like synchronous receive and confirmed delivery may be offered with message system calls, but since applications implement their own shared memory IPC, they would have to provide these services themselves.

## 7. free lists

This was discussed in reading section(s) AD17.2

This was discussed in lecture section(s) 5C,5G

- a. In variable-partition allocation we need to know the size, locations, and neighbors of each chunk. In fixed partition allocation, all of these are constants.
- b. The free list data structures must be designed to optimize:
  - searching for a piece of desired size.
  - breaking a large piece into smaller pieces.
  - coalescing neighbors back together.
- c. we discussed several types of diagnostic information that could be added to free list descriptors and chunks:
  - if we keep allocated memory on a list (as well as free memory) we can audit that list to find memory that has not yet been freed, and perhaps detect

memory leaks.

- address of the allocator (and perhaps time of allocation. This can be recorded at allocation time. If a subsequent audit finds this chunk to be lost, we will know who allocated it (and hence what it was used for).
- we can put pattern-data guard-zones before and after each chunk (at allocation time) and do periodic audits to see that they still contain the correct patterns. This will detect buffer under- or over-run.

## 8. prod/cons w/sems

This was discussed in reading section(s) AD31.4

This was discussed in lecture section(s) 7I

This application probably calls for two different semaphores:

- a work semaphore to allow back-end threads to await requests, The front-end would V the work queue whenever a new request was added to it, and the back-end threads would P the work queue to await work.
- a mutex semaphore to serialize access to the shared queue. All threads (front-end and back-end) would have to P to lock the mutex, and V to release it when adding or removing requests to/from the queue.

The trick is to avoid deadlock (holding one semaphore and then blocking on the other). Nobody holds the mutex while doing a P on the work queue.

**server:**

```
P(mutex)
append to work queue
V(mutex)
V(work queue)
```

**worker:**

```
P(work queue)
P(mutex)
take item off queue
V(mutex)
```

Note that a two-semaphore solution invites deadlock (much like we saw in the semaphore producer/consumer solution we examined in class. I address this by avoiding hold-and-block on the mutex (release the mutex before P'ing the work semaphore).

## 9. Page Fault process

This was discussed in reading section(s) AD21.3-5

This was discussed in lecture section(s) 6C

- a. the trap and low level handling:
  - process reference address that is not yet mapped in
  - CPU generates a page fault exception and traps into the OS
  - first level handler is selected from an in-memory trap vector
  - the PC/PS at time of trap is pushed onto the supervisor mode stack
  - first level handler saves registers and forwards to 2nd level handler.
- b. software loopup, selection, I/O:
  - page fault handler determines that address does indeed refer to a valid, but paged out, page in the process's address space.
  - a free page frame is found, perhaps requiring some other page to be written out
  - I/O request is scheduled to bring in the required page, and we await completion
  - process's page table is adjusted to show location of newly fetched page.
- c. return/resumption:
  - back-up the failed instruction
  - return through the first level handler, which will restore the saved registers.
  - return to usermode with a **return from trap** instruction that will restore the saved PC/PS.
  - resumed process will re-attempt the instruction that had page faulted.

## 10. using Condition Variables

This was discussed in reading section(s) AD30.1

This was discussed in lecture section(s) 7F,7I

- a. The mutex prevents us from missing a wake-up because the condition was signaled, after we checked it, but before we went to sleep.
- b. Sample signal and wait code is:

```

waiter:
    pthread_mutex_lock(&mutex);
    while (!condition)
        pthread_cond_wait(&cv, &mutex);
    pthread_mutex_unlock(&mutex);

signaler:
    pthread_mutex_lock(&mutex);
    condition = True;
    pthread_cond_signal(&cv);
    pthread_mutex_unlock(&mutex);

```

Note that the mutex is held whenever the condition is manipulated or a call is made to either *signal* or *wait*.

- c. The OS will release the mutex after blocking the process (in a call to `pthread_cond_wait`) and reacquire the mutex before returning to the user-mode process.
- d. If the waiter released the mutex prior to calling `pthread_cond_wait`, the signal could be sent before we went to sleep, and we would have missed the wake-up.

## XC. memory allocation mechanisms

This was discussed in reading section(s) AD14

This was discussed in lecture section(s) 5B

Note: this was intended to be *hard question* on this exam, requiring more than mere recollection. Only part (a) was answered in class. Parts (b) and (c) require you to contemplate how the mechanisms might be used.

- a. Stack allocated storage is automatically deallocated when the allocating block exits. Heap storage persists after exiting the block, until it is explicitly freed.
- b. The `sbrk(2)` system call can extend or shrink the data segment but only at its end. We cannot free individually allocated chunks from the middle.
- c. Two likely applications are:
  - allocating very large blocks of memory in their own segments ... where the malloc arena adds little value

- creating multiple malloc arenas (for different clients) each in its own segment).

The first two points were discussed in class. The last point calls for imagination, which is what made this an extra credit problem.